



LOST BOYS CYBER

THREAT INTELLIGENCE • MALWARE ANALYSIS • DETECTION ENGINEERING



Anthropic says US government forced it to disable cybersecurity AI models

Threat Actor Profile

Classification	TLP:CLEAR
Published	2026-06-15
Version	1.0
Author	Lost Boys Cyber
Report Type	Threat Actor Profile

TLP:CLEAR | Lost Boys Cyber | Anthropic says US government forced it to disable cybersecurity AI models - Threat Actor Profile

Published: 2026-06-15 | TLP:CLEAR | Version: 1.0 | Author: Lost Boys Cyber

AI Generation: This report was generated with AI assistance and is provided for informational purposes only. All findings, IOCs, detection rules, hunting queries, attribution assessments, and recommendations must be independently reviewed and validated by a qualified security professional before operational use.

Table of Contents

1. Executive Summary
2. Intelligence Requirement
3. Source Review and Confidence
4. Key Findings
5. Threat Context
6. Detection and Hunting
7. Recommendations
8. References

Anthropic Says U.S. Government Forced It to Disable Cybersecurity AI Models

1. Executive Summary

Anthropic disabled access to Fable 5 and Mythos 5, described in public reporting and the company's statement as its most advanced cybersecurity-capable models, after receiving a U.S. government export-control directive that barred access by foreign nationals. The directive reportedly applied to foreign nationals both outside and inside the United States, including Anthropic employees, which led Anthropic to disable the models for all customers to ensure compliance.

This report treats the event as an AI-enabled cyber capability governance and exposure issue rather than a conventional named threat-actor profile. The relevant actor set includes frontier AI providers, U.S. export-control and national-security authorities, defensive cybersecurity teams using frontier models, and adversaries seeking jailbreaks or unrestricted cyber assistance. The principal defensive implication is that access to cyber-capable AI systems is now a supply-chain and operational-continuity dependency that can change abruptly for legal, national-security, or safety reasons.

Lost Boys Cyber assesses with medium confidence that organizations using frontier AI models for vulnerability research, secure code review, malware triage, or agentic cyber workflows should prepare for model-access volatility, stronger user-nationality and data-retention controls, and increased scrutiny of shadow AI usage. The event does not provide evidence of active exploitation against customer networks, but it does validate a broader risk: high-capability AI models can become both defensive force multipliers and regulated dual-use cyber infrastructure.

Executive Signal	Assessment	Defensive Action
Model access disruption	Fable 5 and Mythos 5 access was abruptly disabled following a U.S. directive	Maintain approved fallback models and offline playbooks for AI-assisted security workflows
Cyber jailbreak concern	Government concern reportedly centered on bypassing safeguards to identify software flaws	Treat prompts, tool-use logs, and model outputs as auditable security telemetry
Regulatory precedent	Reporting frames this as a first-of-kind use of national-security authorities against AI model export/access rather than hardware	Add AI model providers to third-party risk, export-control, and business-continuity reviews
Shadow AI pressure	Users may migrate to alternate or personal tools if approved models disappear	Monitor unsanctioned AI endpoints and enforce approved proxy or CASB controls

2. Intelligence Requirement

Priority Intelligence Requirement	Current Answer	Confidence	Collection Gap
What happened?	Anthropic says it received a U.S. export-control directive requiring suspension of Fable 5 and Mythos 5 access for foreign nationals; it disabled the models for all customers to ensure compliance.	High	The directive itself has not been made public.
Why does this matter to defenders?	The incident demonstrates that frontier AI model access can be disrupted by national-security action, even when the models	Medium	Actual customer impact and restoration timeline remain unclear.

	are used for defensive cyber workflows.		
Is there evidence of adversary exploitation?	Public reporting describes a reported jailbreak concern and vulnerability-identification capability, not confirmed exploitation against customer environments.	Medium	No public technical details or IOCs are available.
Which organizations should act first?	Security teams that depend on frontier models for vulnerability research, secure code review, SOC enrichment, malware triage, or agentic workflow orchestration.	Medium	Internal dependency mapping is required per organization.
What should be monitored?	Approved AI service usage, unsanctioned model endpoints, unusual code/archive uploads to AI platforms, API-key creation, and abrupt drops or errors in sanctioned provider usage.	Medium	Telemetry coverage varies by proxy, EDR, IdP, CASB, and developer platform logging.

Assumptions used in this report:

- The public Anthropic statement and The Record reporting accurately describe the directive at a high level.
- Fable 5 and Mythos 5 are treated as cyber-capable frontier models based on reporting and Anthropic's reference to safeguards for cybersecurity tasks.
- No public exploit chain, malware family, infrastructure, or victim compromise has been disclosed as of this writing.

3. Source Review and Confidence

Source	Contribution	Reliability / Credibility	Analyst Notes
Anthropic statement on Fable 5 and Mythos 5 access	Primary vendor account of the directive, scope, timing, safeguards posture, and disagreement with the government's basis	A2 for Anthropic's own action; B3 for government rationale because the directive is not public	Best source for what Anthropic claims it did and why; not independently sufficient to validate government concerns
The Record by Recorded Future News	Cybersecurity-focused reporting summarizing Anthropic's statement and adding context on export-control precedent and policy dispute	B2	Primary intake source; credible secondary report but dependent on disclosed statements
CyberScoop reporting	Additional reporting on scope, Project Glasswing implications, and industry commentary	B2	Useful corroboration for security-sector impact and dual-use framing
Guardian / Reuters-linked reporting surfaced in search results	Broad corroboration that the action affected access to advanced Anthropic models and was tied to foreign-national restrictions	B2	Used for context only; not relied on for unique technical claims

Lost Boys Cyber analysis	Defensive impact assessment, detection/hunting guidance, and recommended actions	C3	Analytical extrapolation from public facts; confidence is medium because no directive text or technical jailbreak details are public
--------------------------	--	----	--

Confidence statement: Lost Boys Cyber has high confidence that Anthropic publicly reported disabling Fable 5 and Mythos 5 in response to a U.S. government directive. We have medium confidence in the assessed defensive implications because they depend on each organization's AI operating model, telemetry coverage, contractual dependencies, and regulatory exposure. We have low confidence in any claim about the exact jailbreak mechanics or national-security basis because those details are not public.

4. Key Findings

Finding	Confidence	Defensive Implication
The access restriction covered foreign nationals broadly, including Anthropic's own employees according to the vendor statement.	High	AI governance programs need identity, nationality/export-control, and contractual-risk handling that is more mature than generic SaaS access reviews.
Anthropic disabled the models globally for customers to avoid noncompliance with the directive.	High	Model access can fail closed at provider level; SOC, vulnerability-management, and development teams need fallback workflows.
The government's reported concern involved a method of bypassing safeguards or jailbreaking Fable 5 for vulnerability-identification tasks.	Medium	Security teams should preserve logs for prompts, tool calls, and model outputs where AI systems can access code, repositories, or scanners.
Anthropic disputed the severity, saying the observed issues were narrow, previously known, minor, and reproducible with other public models.	Medium	Risk does not disappear when one provider is restricted; adversaries and employees can pivot to other models.
The event may establish a precedent for regulating AI model access as exportable cyber capability, not only regulating chips or hardware.	Medium	Third-party risk reviews should include model capability tier, data-residency, export-control exposure, and sudden-access-loss scenarios.
There are no public network IOCs, malware hashes, domains, exploit payloads, or named adversary infrastructure associated with the incident.	High	Detection should focus on enterprise AI usage, data movement, identity activity, and policy evasion rather than IOC matching.

MITRE ATT&CK mapping is necessarily defensive and behavior-oriented rather than incident-specific:

ATT&CK Tactic	Technique	Relevance
Resource Development	T1588.006 - Obtain Capabilities: Vulnerabilities	Cyber-capable models can assist vulnerability discovery or triage when misused.
Reconnaissance	T1595 - Active Scanning	Agentic AI workflows may combine code analysis, scanning, and triage at scale.
Collection	T1213 - Data from Information	Shadow AI usage may involve uploading

	Repositories	source code, logs, or documents from internal repositories.
Exfiltration	T1567.002 - Exfiltration to Cloud Storage	External AI platforms can become unsanctioned destinations for sensitive code or telemetry.
Defense Evasion	T1562 - Impair Defenses	Jailbreak attempts are not ATT&CK-native but are conceptually analogous to bypassing policy controls around cyber assistance.

5. Threat Context

The incident sits at the intersection of three trends: rapid expansion of agentic cybersecurity AI, government concern over frontier model misuse, and enterprise dependency on externally hosted AI capabilities. Anthropic's statement says the directive arrived at 5:21 p.m. ET, did not provide specific details of the national-security concern, and resulted in the abrupt disablement of Fable 5 and Mythos 5 for all customers. Public reporting further notes that the directive appears unusual because it targeted access to AI models rather than only chips, hardware, or infrastructure.

For defenders, the most important context is not whether this specific jailbreak concern was technically severe. The key issue is that organizations are beginning to embed AI models into workflows that previously relied on locally controlled tooling: secure code review, vulnerability reproduction, exploitability triage, SOC summarization, malware reverse engineering, and autonomous task orchestration. A provider-side access change can therefore degrade a security team's ability to operate, even when no adversary has touched the environment.

Threat actors also benefit from the same ecosystem pressure. If a frontier model is unavailable or heavily restricted, malicious users can pivot to weaker controls, alternate providers, compromised API keys, self-hosted models, or prompt-bypass marketplaces. Defensive monitoring must therefore avoid a single-provider lens. The correct hunting question is not only “Who used Fable 5?” but “Where is sensitive cyber work or internal data leaving approved control planes for any external AI service?”

Relevant exposure categories:

Exposure Category	Example Risk	Likely Telemetry
Approved AI dependency	SOC playbooks or secure-code-review processes fail when model access is revoked	API gateway logs, provider status logs, automation failures, SIEM job errors
Shadow AI substitution	Users upload code or logs to personal AI accounts after enterprise model access changes	Proxy, DNS, CASB, browser history, DLP, EDR network events
Sensitive code disclosure	Repositories, patches, exploit PoCs, or vulnerability details sent to external model endpoints	Git audit logs, upload telemetry, HTTP POST volume, archive creation, endpoint file events
Identity and export-control gap	Access policies cannot distinguish nationality/export-control eligibility, contractor status, or jurisdiction	IdP attributes, HRIS-to-IdP sync, entitlement reviews, audit logs
Adversarial prompt activity	Users attempt jailbreaks or cyber-policy bypasses against approved tools	Prompt logs, model safety logs, application audit tables, API response metadata

6. Detection and Hunting

Detection for this event should be framed as AI-governance telemetry, data-loss prevention, and cyber-workflow resilience. No public IOCs are available. The following content is designed for enterprise adaptation and should be tuned to approved provider lists, proxy schema, and internal AI platform names.

Detection Objective	Signal	Telemetry Source	Priority
Identify sanctioned model disruption	Spikes in 401/403/429/5xx responses or automation failures against Anthropic/Claude endpoints	API gateway, proxy, SaaS connector, application logs	High
Detect shadow AI pivoting	New or unusual access to external AI domains by developers, security staff, or build agents	DNS, proxy, EDR network events, CASB	High
Detect sensitive code movement	Archive creation followed by upload to AI/chat/model endpoints	EDR file + network, DLP, web proxy	High
Monitor prompt-policy abuse	Repeated prompts containing jailbreak, exploit-generation, vulnerability-chain, or safeguard-bypass terms	AI gateway prompt logs, application audit tables	Medium
Review identity/export-control exposure	AI access by users without complete HR/contractor/jurisdiction attributes	IdP, HRIS, SaaS entitlement export	Medium

Sigma rule for proxy or secure web gateway logs:

```

title: External AI Platform Access From Security or Developer Workstations
id: 6ed2fd62-fd80-4bb1-a37e-8a3b222e6b4d
status: experimental
description: Detects access to common external AI service domains from developer or security workstations, useful for identifying shadow AI pivots after sanctioned model access changes.
references:
  - https://therecord.media/anthropic-says-gov-forced-it-to-disable-cyber-ai-models
  - https://www.anthropic.com/news/fable-mythos-access
author: Lost Boys Cyber
date: 2026-06-16
logsource:
  category: proxy
  product: webgateway
detection:
  ai_domains:
    url.domain|contains:
      - anthropic.com
      - claude.ai
      - api.anthropic.com
      - openai.com
      - api.openai.com
      - chatgpt.com
      - gemini.google.com
      - generativelanguage.googleapis.com
      - groq.com
      - mistral.ai
      - openrouter.ai
  sensitive_roles:

```

```

user.department|contains:
- Security
- Engineering
- DevOps
- Research
upload_or_api:
  http.method:
    - POST
    - PUT
  bytes_out|gte: 500000
condition: ai_domains and sensitive_roles and upload_or_api
falsepositives:
- Approved enterprise AI gateway traffic
- Normal developer use of sanctioned coding assistants
level: medium
fields:
- user.name
- src.ip
- url.domain
- http.method
- bytes_out

```

Microsoft Defender / MDE hunt for endpoints contacting AI services after archive creation:

```

let AiDomains = dynamic([
  "anthropic.com", "claude.ai", "api.anthropic.com",
  "openai.com", "api.openai.com", "chatgpt.com",
  "gemini.google.com", "generativelanguage.googleapis.com",
  "groq.com", "mistral.ai", "openrouter.ai"
]);
let ArchiveEvents = DeviceFileEvents
| where Timestamp > ago(14d)
| where FileName endswith ".zip" or FileName endswith ".7z" or FileName endswith ".tar" or FileName
endswith ".gz"
| project DeviceId, DeviceName, InitiatingProcessAccountName, ArchiveTime=Timestamp,
ArchiveName=FileName, FolderPath;
DeviceNetworkEvents
| where Timestamp > ago(14d)
| where RemoteUrl has_any (AiDomains)
| where InitiatingProcessFileName in~ ("chrome.exe", "msedge.exe", "firefox.exe", "curl.exe",
"python.exe", "node.exe", "powershell.exe")
| join kind=leftouter ArchiveEvents on DeviceId
| where isempty(ArchiveTime) or Timestamp between (ArchiveTime .. ArchiveTime + 2h)
| project Timestamp, DeviceName, InitiatingProcessAccountName, InitiatingProcessFileName, RemoteUrl,
RemoteIP, ArchiveName, FolderPath
| order by Timestamp desc

```

Microsoft Sentinel hunt for unusual external AI usage by privileged or security users:

```

let AiTerms = dynamic(["anthropic", "claude", "openai", "chatgpt", "gemini", "mistral", "openrouter",
"groq"]);
let SensitiveUsers = IdentityInfo
| where TimeGenerated > ago(30d)
| where AssignedRoles has_any ("Global Administrator", "Security Administrator", "Privileged Role
Administrator")
  or Department has_any ("Security", "Engineering", "DevOps", "Research")
| summarize arg_max(TimeGenerated, *) by AccountUPN
| project AccountUPN;
CloudAppEvents
| where TimeGenerated > ago(14d)
| where toString(RawEventData) has_any (AiTerms) or Application has_any (AiTerms)
| extend AccountUPN = toString(AccountDisplayName)
| join kind=leftouter SensitiveUsers on AccountUPN
| summarize Events=count(), BytesOut=sum(tolong(RawEventData.BytesSent)), Apps=make_set(Application,
10), ExampleURLs=make_set(ObjectName, 10) by AccountUPN, bin(TimeGenerated, 1d)
| where Events > 10 or BytesOut > 1000000 or isnotempty(AccountUPN1)
| order by TimeGenerated desc

```


Splunk hunt for proxy uploads to external AI endpoints:

```
index=proxy earliest=-14d
(http_method=POST OR http_method=PUT)
(uri_domain="*anthropic.com" OR uri_domain="*claude.ai" OR uri_domain="*openai.com" OR
uri_domain="*chatgpt.com" OR uri_domain="*gemini.google.com" OR uri_domain="*mistral.ai" OR
uri_domain="*openrouter.ai")
| eval mb_out=bytes_out/1024/1024
| stats count sum(mb_out) as total_mb values(uri_domain) as domains values(url) as sample_urls by user
src_ip user_department
| where total_mb > 1 OR count > 20
| sort - total_mb
```

Prompt-audit query pattern for organizations with an internal AI gateway table:

```
AIGatewayPromptLogs
| where TimeGenerated > ago(30d)
| where Provider in~ ("Anthropic", "OpenAI", "Google", "Mistral", "OpenRouter")
| where PromptText has_any ("jailbreak", "bypass", "ignore policy", "exploit chain", "weaponize",
"reverse shell", "CVE", "0day", "vulnerability scanner")
| summarize Prompts=count(), Models=make_set(Model, 10), FirstSeen=min(TimeGenerated),
LastSeen=max(TimeGenerated), Samples=make_set(substring(PromptText, 0, 160), 5) by UserPrincipalName,
AppName
| order by Prompts desc
```

7. Recommendations

Priority	Owner	Recommendation	Success Measure
Critical	CISO / Security Architecture	Inventory security workflows that depend on externally hosted frontier models, including SOC summarization, vulnerability triage, code review, malware analysis, and autonomous agents.	Current AI dependency register includes owner, provider, model tier, data type, fallback path, and business impact.
High	IAM / GRC	Add AI model access to third-party risk and entitlement review processes, including contractor status, jurisdiction/export-control attributes where legally appropriate, and privileged-user access.	Quarterly access review covers AI platforms and approved gateway roles.
High	SOC / Detection Engineering	Deploy monitoring for sanctioned AI outage signals, shadow AI endpoint usage, large uploads to AI services, and prompt-policy abuse.	SIEM contains tuned detections with documented false-positive handling and approved-domain allowlists.
High	Legal / Procurement	Review AI provider terms for model availability, data retention, government-order response, customer notification, and geographic/user eligibility constraints.	Contract register captures provider obligations and disruption escalation contacts.
Medium	Engineering / DevSecOps	Route AI-assisted code review and vulnerability research through approved enterprise	Repository and endpoint monitoring show reduced direct uploads to unsanctioned AI

		gateways instead of direct personal accounts.	platforms.
Medium	Security Operations	Build non-AI and alternate-model fallbacks for critical playbooks, especially incident triage, vuln reproduction, and report generation.	Tabletop exercise demonstrates playbook execution without the primary model provider.
Medium	Privacy / Data Protection	Classify prompts and model outputs as potentially sensitive records when they include source code, logs, vulnerabilities, or customer information.	DLP policy includes AI endpoints and prompt/output retention rules.

Immediate 7-day action plan:

- Confirm whether any internal workflows used Fable 5, Mythos 5, Claude/Anthropic APIs, or Project Glasswing-style access.
- Query proxy and endpoint telemetry for large uploads to AI domains during the 14 days before and after the public disruption.
- Identify users who pivoted from sanctioned model endpoints to personal or unsanctioned AI services.
- Preserve AI gateway logs, prompt logs, and API error logs for at least 30 days while the access issue is active.
- Notify security engineering, legal, and procurement that frontier model availability is now an explicit business-continuity risk.

8. References

- [1] The Record from Recorded Future News. “Anthropic says US government forced it to disable cybersecurity AI models.” Published 2026-06-15. <https://therecord.media/anthropic-says-gov-forced-it-to-disable-cyber-ai-models>
- [2] Anthropic. “Statement on the US government directive to suspend access to Fable 5 and Mythos 5.” Published 2026-06-13. <https://www.anthropic.com/news/fable-mythos-access>
- [3] CyberScoop. “Anthropic disables new models after government calls them a national security concern.” Published 2026-06-13. <https://cyberscoop.com/us-government-anthropic-fable-5-mythos-5-export-controls/>
- [4] The Guardian. “Anthropic to disable its most advanced AI models after US order limiting foreign access.” Published 2026-06-13. <https://www.theguardian.com/technology/2026/jun/13/anthropic-disable-advanced-ai-models-us-government-order>